



Software project: Neural Networks from Scratch


Author: Suvorova Alexandra

Supervisor: Trushin Dmitry, Docent, FCS

Higher School of Economics, 2025

Goal of the project

Implement a fully connected neural network from scratch as a library in C++



Main objectives



Implement a library
with different
activation and loss
functions and
optimizers



Design modular
architecture



Train model using
backpropagation and
optimizers



Evaluate model
performance



Create a clean
interface

Why from Scratch?



Better understanding of core ML principles

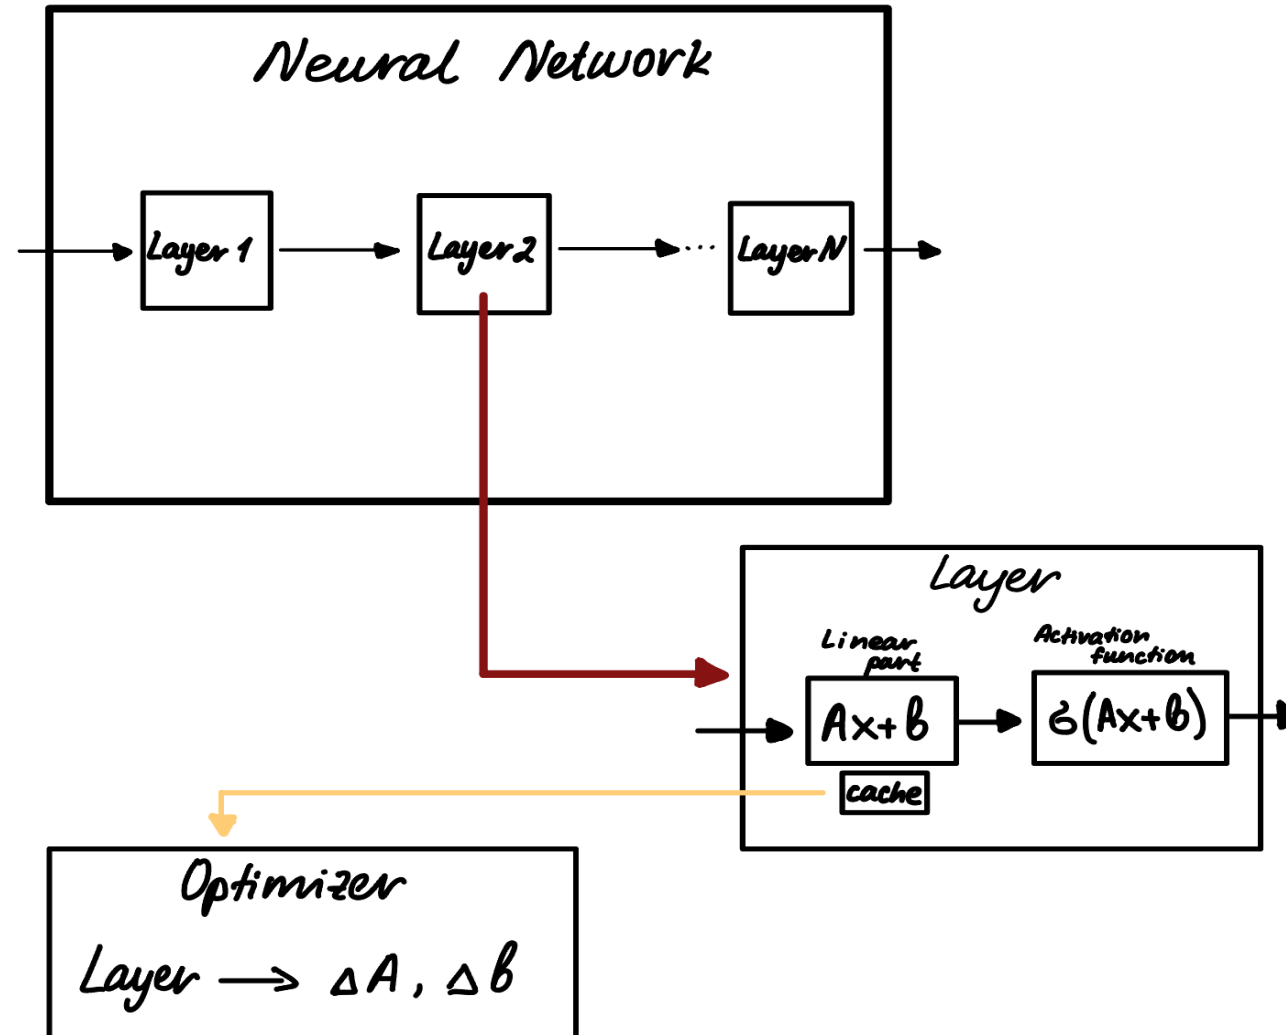


Learn how backpropagation works under the hood



Practice with low-level implementation in C++

Net structure



Tech stack

C++20

Cmake version 3.10

Git as a version control system

LLVM code style

Eigen and EigenRand for linear algebra computations

Features

Fully connected layers

- Linear + Non-linear

Polymorphism

- Type Erasure

Activation Function

- ReLU
- Sigmoid
- Tanh
- Identity
- Softmax

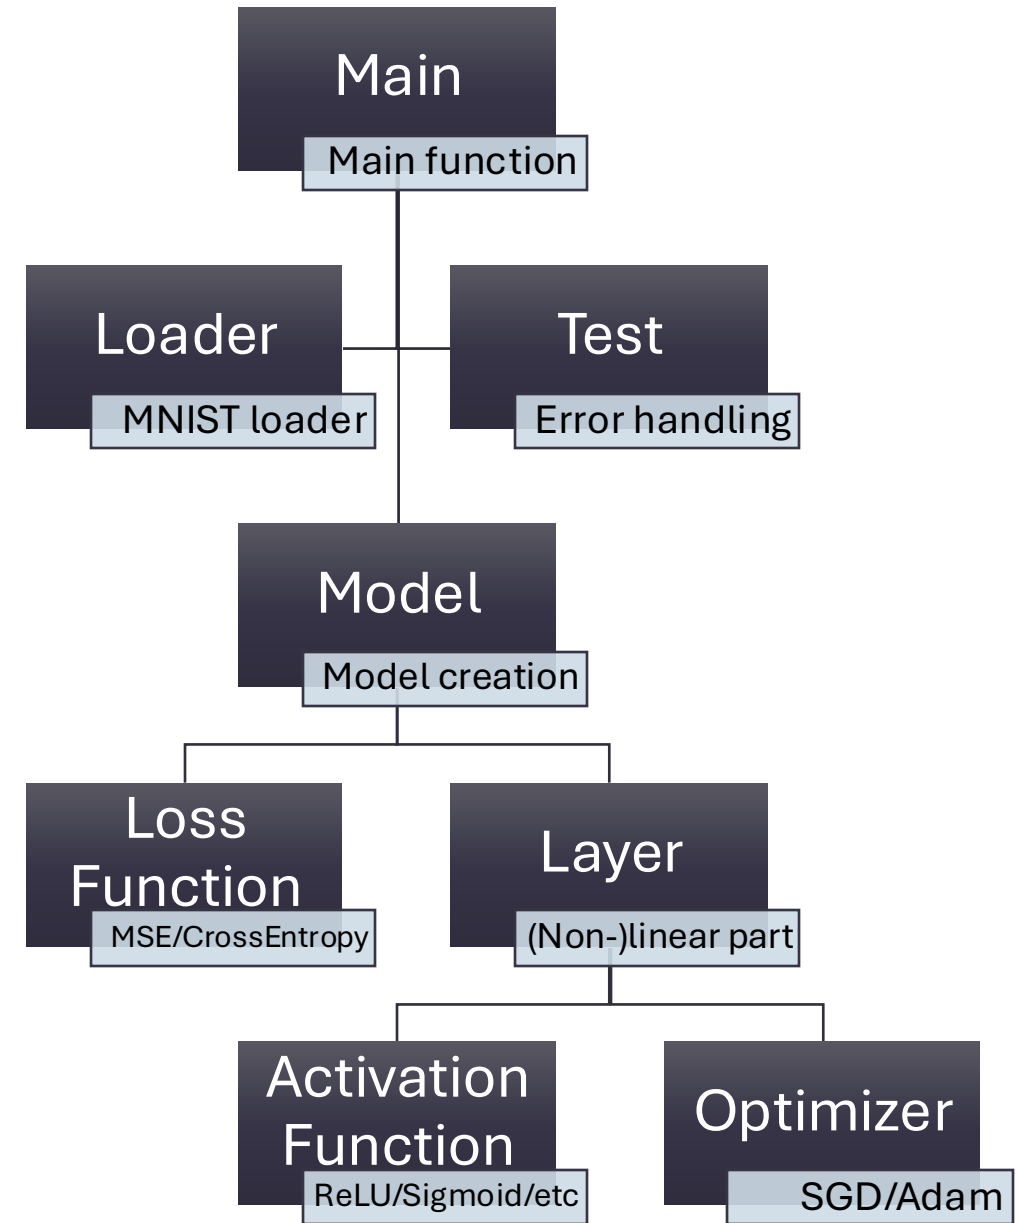
Loss Function

- MSE
- Cross Entropy

Optimizer

- SGD
- Adam

Code Structure and Modules



Code details

- Creation of a Model

```
Optimizer opt = Optimizer::Adam(0.001, 0.9, 0.999, 1e-8);  
Model model({784, 128, 10}, {ActivationFunction::Type::ReLU,  
                               ActivationFunction::Type::Identity});
```

- Launch of training

```
for (int i = 0; i < N; ++i) {  
    model.trainStep(train_images[i], train_targets[i], LossFunction::mseGrad,  
                    opt);  
}
```

Testing

Correctness


- Assert

Accuracy

- $\text{\#successes} / \text{\#all}$

MNIST dataset: handwritten digits

- 60,000 pictures train
- 10,000 pictures test



Confusion matrix

$A[i][j]$ = # of samples of true class i that were predicted as class j





- One epoch of training
- 60,000 pictures – loss computation on each step
- 3 layers – 784-128-10
- ReLU-Identity
- MSE
- Adam




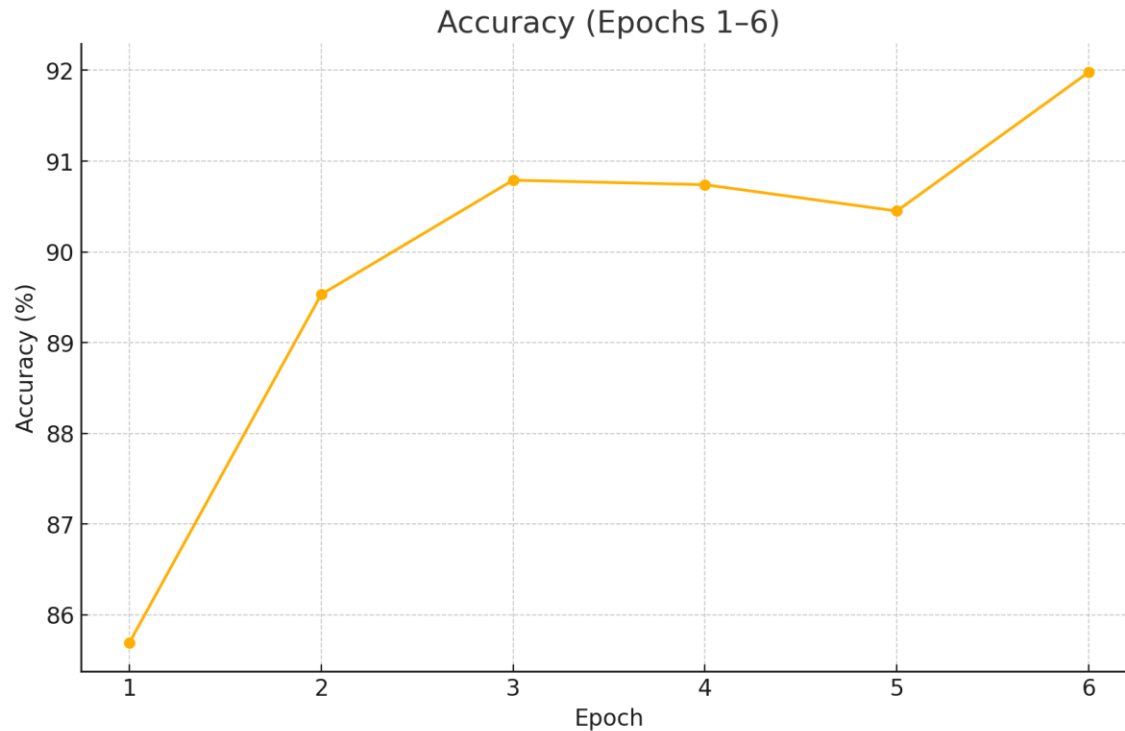
Comparison – 20 epochs

- 5 layers – 784-128-64-32-10
- ReLU-ReLU-ReLU-Softmax
- Cross Entropy
- Adam optimizer



Comparison – 6 epochs

- 5 layers – 784-128-64-32-10
- ReLU-ReLU-ReLU-Softmax
- Cross Entropy
- Adam optimizer



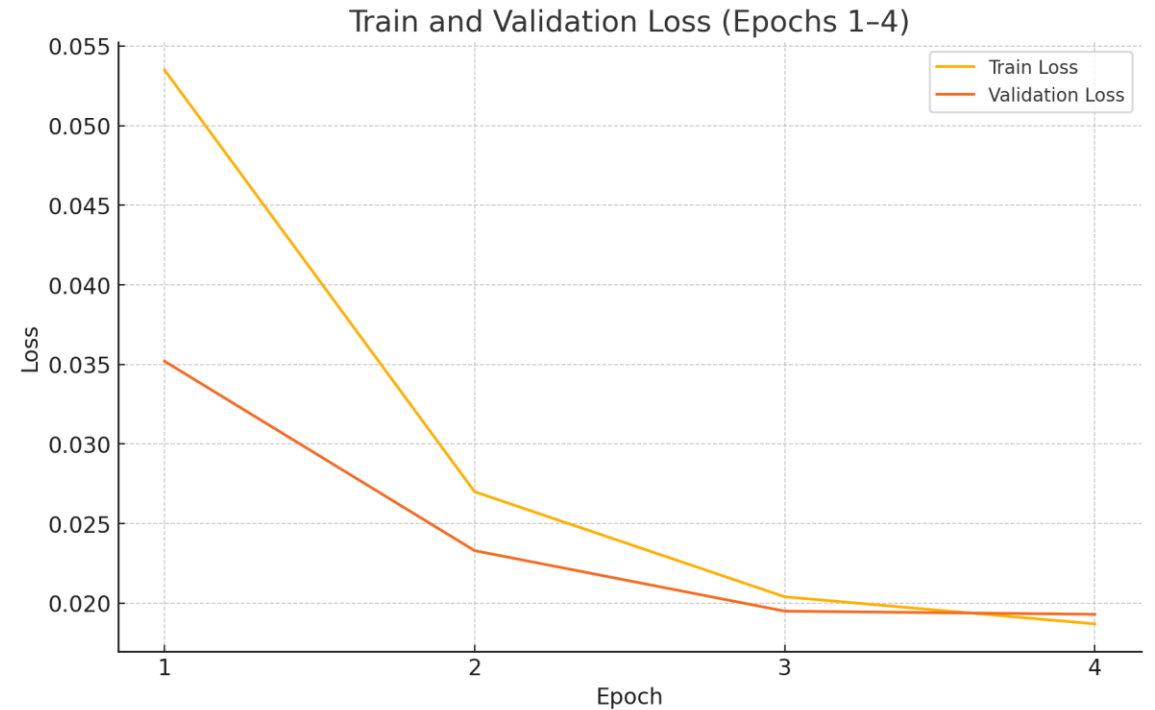

Remarkable result

```
> build git:(dev) x ./neural_net
[OK] All tests passed!
Select model architecture:
1. One hidden layer (ReLU + Identity)
2. Two hidden layers (ReLU + Sigmoid + Identity)
3. Three hidden layers (ReLU + ReLU + ReLU + Softmax)
Enter choice (1/2/3): 2
Enter number of training epochs: 4

=== Training model2 for 4 epoch(s) ===
[=====] 100% (60000/60000) L:0.0535
Epoch 1 finished. Train Loss: 0.0535, Val Loss: 0.0352, Accuracy: 84.8500%
[=====] 100% (60000/60000) L:0.0270
Epoch 2 finished. Train Loss: 0.0270, Val Loss: 0.0233, Accuracy: 89.0200%
[=====] 100% (60000/60000) L:0.0204
Epoch 3 finished. Train Loss: 0.0204, Val Loss: 0.0195, Accuracy: 91.0000%
[=====] 100% (60000/60000) L:0.0187
Epoch 4 finished. Train Loss: 0.0187, Val Loss: 0.0193, Accuracy: 91.0100%
```

Remarkable result

- 4 layers – 784-128-64-10
- ReLU-Sigmoid-Identity
- MSE
- Adam



Future work

01

Integrate GUI

02

Create a manual
for students

03

Test on other
datasets

Summary

Fully connected layers

- Forward and backward propagation
 - Model creation
-

Activation functions:

- ReLU - Tanh - Identity
 - Sigmoid - Softmax
-

Loss functions:

- MSE
 - Cross Entropy
-

Optimizers:

- SGD
 - Adam
-

Tested on a MNIST dataset – 92% accuracy